

Lecture 15

Single-source shortest paths: We saw in Lab 10 that breadth-first search can be used to find the shortest path between two vertices in a graph. But what if the edges in the graph have different lengths? For example, in a graph representing an airport network, edges have associated lengths corresponding to the amount of time it takes to fly from one airport to the next. Then, we might not just be interested in getting from one place to another in as few stops as possible, but we may instead be interested in minimizing total flight time.

A *weighted* graph is a triple (V, E, w) , where w is a *weight* function. Each edge $e \in E$ has a weight $w(e)$, which may be positive, zero, or negative. Now how can we find the shortest path from one vertex to the others in such a graph? This can be done using recursion and memoization. The non-recursive, iterative implementation of this approach is called the *Bellman-Ford* algorithm.

The basic idea is to create a recursive function `shortestPathHelper(x, y, t)` which finds the shortest path from x to y which takes at most t steps. One option is that it is the same as the shortest path taking at most $t - 1$ steps, and the other is that we should travel to some vertex z first in $t - 1$ steps then take the edge (z, y) in the t th step. We recurse on both options and take the better of the two, and we use memoization to make the function faster. Note that if it's possible to get from x to y at all, then it is possible to do so in $n - 1$ steps, where the graph has n vertices, so the length of the shortest path from x to y is `shortestPathHelper(x, y, n-1)`.

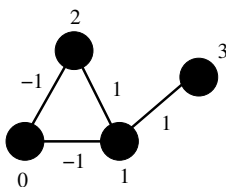


Figure 1: The numbers on edges are weights. This graph has a negative weight cycle $0 \rightarrow 1 \rightarrow 2 \rightarrow 0$.

In some graphs though, such as those in Figure 1, there is no shortest path from some vertex to another. For example, to go from vertex 0 to 3, we could take the route $0 \rightarrow 1 \rightarrow 3$ for a total length of $-1 + 1 = 0$. However note that the cycle $0 \rightarrow 1 \rightarrow 2 \rightarrow 0$ has a total length of -1 . Thus, by repeatedly going on this cycle over and over again, we can make our total length arbitrarily small before finally heading over to vertex 3. Thus, in essence, the length of the shortest path from 0 to 3 is $-\infty$. We modify our `shortestPath` code to detect such negative-weight cycles. If any such cycle is found, starting from our starting vertex x , then we return -1 . Otherwise we return a list with all the shortest path distances.

How can we detect a negative-weight cycle? Let such a reachable cycle be v_0, v_2, \dots, v_{k-1} . Let $d[u]$ be the shortest path distance from x to u taking at most $n - 1$ steps. For the sake of contradiction, assume that we cannot improve the distance to any of the v_i by looking at paths of length n . That means that $d[v_i] \leq d[v_{i-1}] + w(v_{i-1}, v_i)$ for all i (with the understanding that v_{-1} is just v_{k-1}). Summing up all these inequalities gives

$$\sum_{i=1}^k d[v_i] \leq \sum_{i=1}^k d[v_{i-1}] + \sum_{i=1}^k w(v_{i-1}, v_i)$$

Since each $d[v_i]$ appears exactly once in the summations on both sides, we can cancel to then get

$$0 \leq \sum_{i=1}^k w(v_{i-1}, v_i).$$

This is a contradiction, since we assumed that this cycle had negative total weight.

Our implementation now follows.

```
# returns length of shortest path from x to y using at most t steps
def shortestPathHelper(B, x, y, t, mem, seen):
    if t == 0:
        if x == y:
            return 0
        return float('infinity')
    elif seen[y][t]:
        return mem[y][t]

    seen[y][t] = True

    # first option: do it in t-1 steps
    ans = shortestPathHelper(B, x, y, t-1, mem, seen)

    # second option: go to a vertex z that has an edge to y first, in
    # at most t-1 steps, then take the edge (z, y)
    for p in B[y]:
        z = p[0]
        weight = p[1]
        val = shortestPathHelper(B, x, z, t-1, mem, seen)
        ans = min(ans, weight + val)

    mem[y][t] = ans
    return ans

# A is the adjacency list of the graph
# A[u][i][0] is the ith neighbor of vertex u, and A[u][i][1] is the
# weight of the edge (u, A[u][i][0])
#
# returns a list L so that L[j] is the length of the shortest path
# from x to j, assuming no negative-weight cycle is reachable from
# i. returns -1 if a negative-weight cycle is reachable from i.
def shortestPath(A, x):
    mem = []
    # mem[i][j] is float('infinity') if we can't get from x to i in at
    # most j steps. Otherwise, it's the length of the shortest path
    # from x to i taking at most j steps.
```

```

for i in xrange(len(A)):
    mem += [[float('infinity')]*(len(A)+1)]

seen = []
# seen[i][j] is True if we've already filled in mem[i][j] and is
# False otherwise
for i in xrange(len(A)):
    seen += [[False]*(len(A)+1)]

# B is an inverse adjacency list. B[i] is a list of all vertices
# j such that (j, i) is an edge, plus the weight of the edge
B = []
for i in xrange(len(A)):
    B += [[]]
for i in xrange(len(A)):
    for p in A[i]:
        # p is the pair [j, weight(i, j)]
        B[p[0]] += [[i, p[1]]]

# check if a negative weight cycle is reachable from x
for z in xrange(len(A)):
    val1 = shortestPathHelper(B, x, z, len(A) - 1, mem, seen)
    val2 = shortestPathHelper(B, x, z, len(A), mem, seen)
    if val2 < val1:
        return -1

L = []
for y in xrange(len(A)):
    L += [shortestPathHelper(B, x, y, len(A) - 1, mem, seen)]
return L

```

Naïvely one could say that the running time of the algorithm is $O(n^3)$: in the memoized helper function there are n choices for y , n choices for t , and the loop over $B[y]$ might run for $n - 1$ steps. For some graphs this could happen, but in fact the algorithm's running time is $\Theta(n(m + n))$, where the graph has m edges. nm is always at most n^3 since m is at most n^2 , but it can be a lot faster if the graph doesn't have too many edges. The reason the running time is $\Theta(n(m + n))$ is the following. Look at the for loop “for p in B[y]” in `shortestPathHelper`. The total number of (p, y) values this loop executes with is exactly m : each (p, y) pair corresponds to an edge in the graph. Then, there are n possible values of t , giving nm .

The n^2 term comes from there being n possible values of both y and t , though this term can be removed with a better implementation, which we give below. The below implementation is an iterative implementation of the approach above, but written iteratively instead of recursively. It is known as the Bellman-Ford algorithm. The code is also quite a bit shorter than the recursive implementation given above.

```

def bellmanFord(A, x):
    # E is a list of edges with weights
    E = []
    for i in xrange(len(A)):
        for p in A[i]:
            E += [[i] + p]
    # dist[i] is the length of the shortest path to i
    dist = [float('infinity')]*len(A)
    dist[x] = 0
    for i in xrange(len(A) - 1):
        for e in E:
            u = e[0]
            v = e[1]
            weight = e[2]
            dist[v] = min(dist[v], dist[u] + weight)

    # look for negative weight cycles
    for e in E:
        u = e[0]
        v = e[1]
        weight = e[2]
        if dist[u] + weight < dist[v]:
            return -1

    L = []
    for i in xrange(len(A)):
        L += [dist[i]]
    return L

```

All pairs shortest paths: What about finding all shortest path lengths between all pairs of vertices? Here we use memoization yet again. Let's assume the graph has no negative weight cycles. An iterative, non-recursive version of the below approach is known as the *Floyd-Warshall* algorithm.

The approach is to let `shortestPathHelper(u, v, k)` be the length of the shortest path where the intermediate vertices come from the set $\{0, 1, \dots, k-1\}$. For such a shortest path we have two choices: either don't use vertex $k-1$ at all, so that the answer is `shortestPathHelper(u, v, k-1)`, or use it so that the answer is `shortestPathHelper(u, k-1, k-1) + shortestPathHelper(k-1, v, k-1)`. When $k = 0$ we cannot use intermediate vertices, which means the answer is just the weight of the edge from u to v .

```

def shortestPathHelper(u, v, k, w, mem, seen):
    if k == 0:
        return w[u][v]
    elif seen[u][v][k]:
        return mem[u][v][k]

```

```

seen[u][v][k] = True

option1 = shortestPathHelper(u, v, k-1, w, mem, seen)
val1 = shortestPathHelper(u, k-1, k-1, w, mem, seen)
val2 = shortestPathHelper(k-1, v, k-1, w, mem, seen)
mem[u][v][k] = min(option1, val1 + val2)
return mem[u][v][k]

# w is a matrix of edge weights, i.e. w[i][j] is the weight of the
# edge (i,j). If that edge doesn't exist, we assume w[i][j] is then
# float('infinity')
def shortestPath(w):
    ans = []
    for i in xrange(len(w)):
        ans += [[-1]*len(w)]

    mem = []
    for i in xrange(len(w)):
        l = []
        for j in xrange(len(w)):
            l += [[-1]*(len(w)+1)]
        mem += [l]

    # seen[i][j][k] is True if we've already filled in mem[i][j][k],
    # and it is False otherwise
    seen = []
    for i in xrange(len(w)):
        l = []
        for j in xrange(len(w)):
            l += [[False]*(len(w)+1)]
        seen += [l]

    for i in xrange(len(w)):
        for j in xrange(len(w)):
            ans[i][j] = shortestPathHelper(i, j, len(w), w, mem, seen)
    return ans

```

In fact, a non-recursive implementation of this approach is much slicker and amounts to just three for loops. This is known as the Floyd-Warshall algorithm. It also uses less memory than the implementation above ($\Theta(n^2)$ instead of $\Theta(n^3)$).

```

def floydWarshall(w):
    # now dist is a copy of the weight function
    dist = w[:]

    for k in xrange(len(w)):
        for u in xrange(len(w)):
            for v in xrange(len(w)):
                dist[u][v] = min(dist[u][v], dist[u][k] + dist[k][v])
    return dist

```

Both the iterative and recursive implementations of Floyd-Warshall above take $\Theta(n^3)$ time, though the iterative implementation requires less memory.